

Course: MScCS

Year: 2020-2022

Semester: IV

Program: Computer Science

Subject: Big Data Engineering Tools and Frameworks

Subject Code: PS-SCS-404

Seat Number: KSMSCCS012

Name: Tanay Kulkarni

University: HSNC University

College: KC College, Churchgate

Signature

Date

Practical No 1

Steps for Install Hadoop on Windows Based Platform

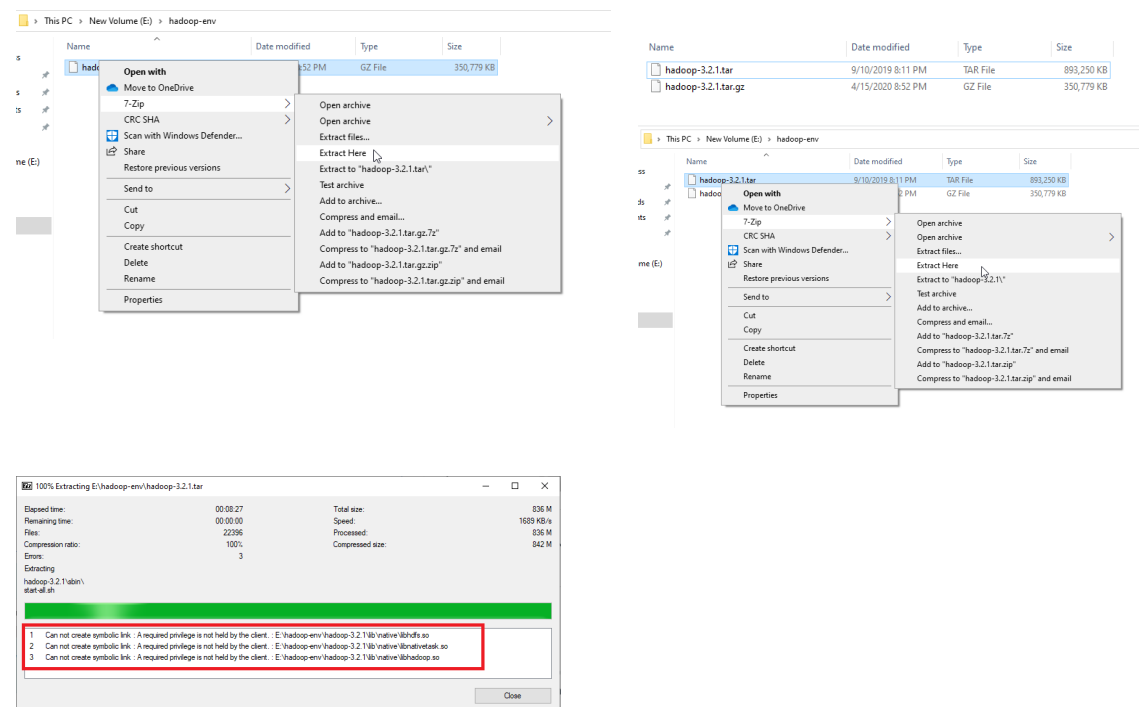
First, we need to make sure that the following prerequisites are installed:

1. Java 8 runtime environment (JRE): Hadoop 3 requires a Java 8 installation. I prefer using the offline installer.
2. Java 8 development Kit (JDK)
3. To unzip downloaded Hadoop binaries, we should install 7zip.
4. I will create a folder "E:\hadoop-env" on my local machine to store downloaded files.

2. Download Hadoop binaries

The first step is to download Hadoop binaries from the official website.

<https://www.apache.org/dyn/closer.cgi/hadoop/common/hadoop-3.2.1/hadoop-3.2.1.tar.gz>



After unpacking the package, add the Hadoop native IO libraries, which can be found in the following GitHub repository:


<https://github.com/cdarlint/winutils>

Since we are installing Hadoop 3.2.1, download the files located in

<https://github.com/cdarlint/winutils/tree/master/hadoop-3.2.1/bin> and copy them into the "hadoop-3.2.1\bin" directory.

3. Setting up environment variables

After installing Hadoop and its prerequisites, we should configure the environment variables to define Hadoop and Java default paths.



```
Windows PowerShell
Copyright (C) Microsoft Corporation. All rights reserved.

Try the new cross-platform PowerShell https://aka.ms/pscore6

PS C:\Users\HFadl> hadoop -version
java version "1.8.0_251"
Java(TM) SE Runtime Environment (build 1.8.0_251-b08)
Java HotSpot(TM) 64-Bit Server VM (build 25.251-b08, mixed mode)
PS C:\Users\HFadl>
```

open "hdfs-site.xml" file located in "%HADOOP_HOME%\etc\hadoop" directory, and we should add the following properties within the <configuration></configuration> element:

```
<property>
<name>dfs.replication</name>
<value>1</value>
</property>
<property>
<name>dfs.namenode.name.dir</name>
<value>file:///E:/hadoop-env/hadoop-3.2.1/data/dfs/namenode</value>
</property>
<property>
<name>dfs.datanode.data.dir</name>
<value>file:///E:/hadoop-env/hadoop-3.2.1/data/dfs/datanode</value>
</property>
```

configure the name node URL adding the following XML code into the <configuration></configuration> element within "core-site.xml":

```
<property>
<name>fs.default.name</name>
<value>hdfs://localhost:9820</value>
</property>
```

add the following XML code into the <configuration></configuration> element within "mapred-site.xml":

```
<property>
<name>mapreduce.framework.name</name>
<value>yarn</value>
<description>MapReduce framework name</description>
</property>
```

add the following XML code into the <configuration></configuration> element within "yarn-site.xml":

```
<property>
<name>yarn.nodemanager.aux-services</name>
<value>mapreduce_shuffle</value>
```


<description>Yarn Node Manager Aux Service</description>
</property>

hdfs namenode -format

```
Administrator: Windows PowerShell

PS E:\hadoop-env\hadoop-3.2.1\sbin> .\start-dfs.cmd
PS E:\hadoop-env\hadoop-3.2.1\sbin>
```

```
PS E:\hadoop-env\hadoop-3.2.1\sbin> jps
14560 DataNode
4960 ResourceManager
5936 NameNode
768 NodeManager
14636 Jps
PS E:\hadoop-env\hadoop-3.2.1\sbin>
```



Hadoop

Overview

Datanodes

Snapshot

Startup Progress

Utilities

Datanode Information

In operation

Node	Last contact	Admin State	Capacity	Used	Non DFS Used	Remaining	Blocks	Block pool used	Failed Volumes	Version
vee-desktop (127.0.0.1:50010)	0	In Service	91.54 GB	612 KB	23.57 GB	67.97 GB	17	612 KB (0%)	0	2.3.0

Decommissioning

Node	Last contact	Under replicated blocks	Blocks with no live replicas	Under Replicated Blocks in files under construction
------	--------------	-------------------------	------------------------------	---

Hadoop, 2014.

LEGEND

UI

All Applications

Cluster

About

Nodes

Node Labels

Applications

NEW

SAVING

SUBMITTED

ACCEPTE

RUNNING

FINISHED

FAILED

QUEUED

Scheduler

Tools

Cluster Metrics

Apps Submitted

Apps Pending

Apps Running

Apps Completed

Containers Running

Memory Used

Memory Total

Memory Reserved

VCores

Cluster Nodes Metrics

Active Nodes

Decommissioning Nodes

Decommissioned Nodes

Lost Nodes

Unhealthy Nodes

Scheduler Metrics

Scheduler Type

Scheduling Resource Type

Minimum Allocation

Maximum Allocation

Capacity Scheduler

memory-m8 (jvm8g1, vcores)

memory-1024_vcores 1x

memory-5102_vcores 4x

0

Show 25 entries

ID

User

Name

Application

Queue

Priority

Start Time

Launch Time

Finish Time

State

Final Status

Running Container

Allocated CPU

Allocated V-Cores

Reserved Memory MB

Reserved V-Cores

Showing 0 to 0 of 0 entries

No data available in table

Practical No 2

Hadoop Word Count

Code:

```
import java.io.IOException;
import java.util.*;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.conf.*;
import org.apache.hadoop.io.*;
import org.apache.hadoop.mapreduce.*;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.input.TextInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
import org.apache.hadoop.mapreduce.lib.output.TextOutputFormat;
public class WordCount {
    public static class Map extends Mapper<LongWritable, Text, Text, IntWritable> {
        private final static IntWritable one = new IntWritable(1);
        private Text word = new Text();
        public void map(LongWritable key, Text value, Context context) throws IOException,
        InterruptedException {
            String line = value.toString();
            StringTokenizer tokenizer = new StringTokenizer(line);
            while (tokenizer.hasMoreTokens()) {
                word.set(tokenizer.nextToken());
                context.write(word, one);
            }
        }
    }

    public static class Reduce extends Reducer<Text, IntWritable, Text, IntWritable> {

        public void reduce(Text key, Iterable<IntWritable> values, Context context)
        throws IOException, InterruptedException {
            int sum = 0;
            for (IntWritable val : values) {
                sum += val.get();
            }
            context.write(key, new IntWritable(sum));
        }
    }

    public static void main(String[] args) throws Exception {
        Configuration conf = new Configuration();
        conf.set("mapred.job.tracker", "hdfs://localhost:50001");
        conf.set("fs.default.name", "hdfs://localhost:50000");
        Job job = new Job(conf, "wordcount");

        job.setJarByClass(WordCount.class);
        job.setOutputKeyClass(Text.class);
        job.setOutputValueClass(IntWritable.class);
    }
}
```

```

job.setMapperClass(Map.class);
job.setReducerClass(Reduce.class);

job.setInputFormatClass(TextInputFormat.class);
job.setOutputFormatClass(TextOutputFormat.class);

FileInputFormat.addInputPath(job, new Path(args[0]));
FileOutputFormat.setOutputPath(job, new Path(args[1]));

job.waitForCompletion(true);
}
}

```

Output:

```

Administrator: Command Prompt
FILE: Number of large read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=1999
HDFS: Number of bytes written=120
HDFS: Number of read operations=6
HDFS: Number of large read operations=0
HDFS: Number of write operations=2
Job Counters
  Launched map tasks=1
  Launched reduce tasks=1
  Data-local map tasks=1
  Total time spent by all maps in occupied slots (ms)=2180
  Total time spent by all reduces in occupied slots (ms)=2442
  Total time spent by all map tasks (ms)=2180
  Total time spent by all reduce tasks (ms)=2442
  Total vcore-milliseconds taken by all map tasks=2180
  Total vcore-milliseconds taken by all reduce tasks=2442
  Total megabyte-milliseconds taken by all map tasks=2232320
  Total megabyte-milliseconds taken by all reduce tasks=2500608
Map-Reduce Framework
  Map input records=30
  Map output records=390
  Map output bytes=2730
  Map output materialized bytes=195
  Input split bytes=111
  Combine input records=390
  Combine output records=21
  Reduce input groups=21
  Reduce shuffle bytes=195
  Reduce input records=21
  Reduce output records=21
  Spilled Records=42
  Shuffled Maps =1
  Failed Shuffles=0
  Merged Map outputs=1
  GC time elapsed (ms)=70
  CPU time spent (ms)=764
  Physical memory (bytes) snapshot=471478272
  Virtual memory (bytes) snapshot=619429888
  Total committed heap usage (bytes)=353894400
Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0
File Input Format Counters
  Bytes Read=1888
File Output Format Counters
  Bytes Written=120
C:\>

```

```

Apache 1
Foundation 1
Software 1
The 1
This 1
by 1
developed 1
includes 1
product 1
software 1

```

Practical No 3

Example working with Hadoop Map Reduce

Given below is the data regarding the electrical consumption of an organization. It contains the monthly electrical consumption and the annual average for various years.

If the above data is given as input, write applications to process it and produce results such as finding the year of maximum usage, year of minimum usage, and so on.

	Jan	Feb	Mar	Apr	May	Jun	Jul	Aug	Sep	Oct	Nov	Dec	Avg
1979	23	23	2	43	24	25	26	26	26	26	25	26	25
1980	26	27	28	28	28	30	31	31	31	30	30	30	29
1981	31	32	32	32	33	34	35	36	36	34	34	34	34
1984	39	38	39	39	39	41	42	43	40	39	38	38	40
1985	38	39	39	39	39	41	41	41	00	40	39	39	45

```
package hadoop;
```

```
import java.util.*;
```

```
import java.io.IOException;
```

```
import java.io.IOException;
```

```
import org.apache.hadoop.fs.Path;
```

```
import org.apache.hadoop.conf.*;
```

```
import org.apache.hadoop.io.*;
```

```
import org.apache.hadoop.mapred.*;
```

```
import org.apache.hadoop.util.*;
```

```
public class ProcessUnits {
```

```
    //Mapper class
```

```
    public static class E_EMapper extends MapReduceBase implements
```

```
Mapper<LongWritable ,/*Input key Type */
```

```
Text,          /*Input value Type*/
```

```
Text,          /*Output key Type*/
```

```
IntWritable>    /*Output value Type*/
```

```
{
```

```
    //Map function
```

```
    public void map(LongWritable key, Text value,
```

```
OutputCollector<Text, IntWritable> output,
```

```
Reporter reporter) throws IOException {
```

```
    String line = value.toString();
```

```
    String lasttoken = null;
```

```
    StringTokenizer s = new StringTokenizer(line,"\\t");
```

```
    String year = s.nextToken();
```

```
    while(s.hasMoreTokens()) {
```

```
        lasttoken = s.nextToken();
```

```

    }
    int avgprice = Integer.parseInt(lasttoken);
    output.collect(new Text(year), new IntWritable(avgprice));
  }
}

//Reducer class
public static class E_EReduce extends MapReduceBase implements Reducer< Text, IntWritable,
Text, IntWritable > {

    //Reduce function
    public void reduce( Text key, Iterator <IntWritable> values,
    OutputCollector<Text, IntWritable> output, Reporter reporter) throws IOException {
        int maxavg = 30;
        int val = Integer.MIN_VALUE;

        while (values.hasNext()) {
            if((val = values.next().get())>maxavg) {
                output.collect(key, new IntWritable(val));
            }
        }
    }
}

//Main function
public static void main(String args[])throws Exception {
    JobConf conf = new JobConf(ProcessUnits.class);

    conf.setJobName("max_eletricityunits");
    conf.setOutputKeyClass(Text.class);
    conf.setOutputValueClass(IntWritable.class);
    conf.setMapperClass(E_EMapper.class);
    conf.setCombinerClass(E_EReduce.class);
    conf.setReducerClass(E_EReduce.class);
    conf.setInputFormat(TextInputFormat.class);
    conf.setOutputFormat(TextOutputFormat.class);

    FileInputFormat.setInputPaths(conf, new Path(args[0]));
    FileOutputFormat.setOutputPath(conf, new Path(args[1]));

    JobClient.runJob(conf);
}
}

```

\$ mkdir units

\$ javac -classpath hadoop-core-1.2.1.jar -d units ProcessUnits.java

\$ jar -cvf units.jar -C units/ .

\$HADOOP_HOME/bin/hadoop fs -mkdir input_dir

\$HADOOP_HOME/bin/hadoop fs -put /home/hadoop/sample.txt input_dir

\$HADOOP_HOME/bin/hadoop fs -ls input_dir/

\$HADOOP_HOME/bin/hadoop jar units.jar hadoop.ProcessUnits input_dir output_dir

Map-Reduce Framework

```
Map input records = 5
Map output records = 5
Map output bytes = 45
Map output materialized bytes = 67
Input split bytes = 208
Combine input records = 5
Combine output records = 5
Reduce input groups = 5
Reduce shuffle bytes = 6
Reduce input records = 5
Reduce output records = 5
Spilled Records = 10
Shuffled Maps = 2
Failed Shuffles = 0
Merged Map outputs = 2
GC time elapsed (ms) = 948
CPU time spent (ms) = 5160
Physical memory (bytes) snapshot = 47749120
Virtual memory (bytes) snapshot = 2899349504
Total committed heap usage (bytes) = 277684224
```

File Output Format Counters

```
Bytes Written = 40
```

```
$HADOOP_HOME/bin/hadoop fs -ls output_dir/
```

```
$HADOOP_HOME/bin/hadoop fs -cat output_dir/part-00000
```

```
1981 34
1984 40
1985 45
```

copy the output folder from HDFS to the local file system.

```
$HADOOP_HOME/bin/hadoop fs -cat output_dir/part-00000/bin/hadoop dfs get output_dir
/home/hadoop
```

Practical No 4

a. Write a Scala program to print "hello world"

```
object MainObject
{
  def main(args:Array[String])
  {
    print("Hello World")
  }
}
```

```
Hello World
warning: 1 deprecation (s
```

b. Write a Scala program to compute the sum of the two given integer value, if the value are the same then return their sum.

```
object MainObject
{
  def main(args:Array[String])
  {
    val a = scala.io.StdIn.readInt()
    val b = scala.io.StdIn.readInt()
    if(a==b){
      print(s"Sum of $a and $b is " + (a + b))
    }else {
      print(s"$a and $b are not same")
    }
  }
}
```

```
Sum of 32 and 32 is 64
warning: 1 deprecation (s
```

c. Write a Scala program to get the absolute difference between n to 51. If n is greater than 51 , error message should display.

```
object MainObject
{
  def main(args:Array[String])
  {
    val n = scala.io.StdIn.readInt()
    val n_abs = n.abs
    if(n_abs > 51) {
      println(s"Invalid Input")
    }else {
      val diff = 51 - n_abs
      println(s"absolute difference between $n and 51 is $diff")
    }
  }
}
```

```
absolute difference between 32 and 51 is 19
```

Practical No 5

a. Write a Scala program to check if a given number is present in first or the last position of given array

object MainObject

```
{
  def main(args:Array[String])
  {
    val list = List(4,3,5,66,8,3,2,1,9,8)
    val n = scala.io.StdIn.readInt()
    if(n == list.head){
      println(s"$n is first element of $list")
    }else if(n == list.last){
      println(s"$n is last element of $list")
    }else {
      println(s"$n is not first or last element of list")
    }
  }
}
```

```
8 is last element of List(4, 3, 5, 66, 8, 3, 2, 1, 9, 8)
```

b. Write a Scala program to find the maximum and minimum value of an array of integers.

object MainObject

```
{
  def main(args:Array[String])
  {
    val list =
List(4,3,5,66,8,3,2,1,9,8)
    val mx = list.max
    val mn = list.min
    println(s"max element of list $list is $mx and min element is $mn")
  }
}
```

```
max element of list List(4, 3, 5, 66, 8, 3, 2, 1, 9, 8) is 66 and min element is 1
```

c. Write a Scala program to find the common element between two arrays of string.

object MainObject

```
{
  def main(args:Array[String])
  {
    val s1 = List("MiniGunner", "Archer", "ElectroShocker", "Ranger")
    val s2 = List("MiniGunner", "AcePilot", "MilitaryBase", "Ranger")
    val common = s1.intersect(s2)
    println(s"common elements between $s1 and $s2 are $common")
  }
}
```

```
common elements between
List(MiniGunner, Archer, ElectroShocker, Ranger) and
List(MiniGunner, AcePilot, MilitaryBase, Ranger) are
List(MiniGunner, Ranger)
```

Practical No 6

a. Write a Scala program to calculate the length of a given list.

```
object MainObject
```

```
{
  def main(args:Array[String])
  {
    val s1 = List("MiniGunner", "Archer", "ElectroShocker", "Ranger")
    val len = s1.length
    println(s"List of length $s1 is $len")
  }
}
```

```
List of length List(MiniGunner, Archer, ElectroShocker, Ranger) is 4
```

b. Write a Scala program to check a given list is a palindrome or not.

```
object MainObject
```

```
{
  def main(args:Array[String])
  {
    val str = scala.io.StdIn.readLine()
    val str_r = str.reverse
    if(str == str_r){
      println(s"$str is a pallindrom")
    }else {
      println(s"$str is not a pallindrom")
    }
  }
}
```

```
HelleselleH is a pallindrom
```

c. Write a Scala program to reverse a given list.

```
object MainObject
```

```
{
  def main(args:Array[String])
  {
    val str = scala.io.StdIn.readLine()
    val str_r = str.reverse
    if(str == str_r){
      println(s"$str is a pallindrom")
    }else {
      println(s"$str is not a pallindrom")
    }
  }
}
```

```
list is: List(1, 2, 4, 5, 234, 23423)
reverse is: List(23423, 234, 5, 4, 2, 1)
```

Practical No 7

Working with Spark

- 1.What are the column name
- 2.What does the Schema look like
- 3.Print out the first 5 columns
- 4.Display mean, count, stdev, min, max
5. Create a new dataframe with a column called hvratio that is the ratio of the high price versus volume of stock traded for a day

```
from pyspark.sql import SparkSession
spark = SparkSession.builder.appName('walmart').getOrCreate()
df = spark.read.csv('walmart_stock.csv', inferSchema=True, header=True)
print(df.columns)
df.printSchema()
for line in df.head(5):
    print(line, '\n')
df.describe().show()
df_hv = df.withColumn('HV Ratio', df['High']/df['Volume']).select(['HV
Ratio'])
df_hv.show()
```

```
['Date', 'Open', 'High', 'Low', 'Close', 'Volume', 'Adj Close']
root
|-- Date: timestamp (nullable = true)
|-- Open: double (nullable = true)
|-- High: double (nullable = true)
|-- Low: double (nullable = true)
|-- Close: double (nullable = true)
|-- Volume: integer (nullable = true)
|-- Adj Close: double (nullable = true)
```

```
Row(Date=datetime.datetime(2012, 1, 3, 0, 0), Open=59.970001, High=61.060001, Low=59.869999, Close=60.330002, Volume=1
Row(Date=datetime.datetime(2012, 1, 4, 0, 0), Open=60.209998999999996, High=60.349998, Low=59.470001, Close=59.7099989
Row(Date=datetime.datetime(2012, 1, 5, 0, 0), Open=59.349998, High=59.619999, Low=58.369999, Close=59.419998, Volume=1
Row(Date=datetime.datetime(2012, 1, 6, 0, 0), Open=59.419998, High=59.450001, Low=58.869999, Close=59.0, Volume=806940
Row(Date=datetime.datetime(2012, 1, 9, 0, 0), Open=59.029999, High=59.549999, Low=58.919998, Close=59.18, Volume=66793
```

summary	Open	High	Low	Close	Volume	Adj Close
count	1258	1258	1258	1258	1258	1258
mean	72.35785375357709	72.83938807631165	71.9186009594594	72.38844998012726	8222093.481717011	67.23883848728146
stddev	6.76809024470826	6.768186808159218	6.744075756255496	6.756859163732991	4519780.8431556	6.722609449996857
min	56.389998999999996	57.060001	56.299999	56.419998	2094900	50.363689
max	90.800003	90.970001	89.25	90.470001	80898100	84.91421600000001

```

+-----+
|              HV Ratio |
+-----+
| 4.819714653321546E-6 |
| 6.290848613094555E-6 |
| 4.669412994783916E-6 |
| 7.367338463826307E-6 |
| 8.915604778943901E-6 |
| 8.644477436914568E-6 |
| 9.351828421515645E-6 |
| 8.29141562102703E-6 |
| 7.712212102001476E-6 |
| 7.071764823529412E-6 |
| 1.015495466386981E-5 |
| 6.576354146362592... |
| 5.90145296180676E-6 |
| 8.547679455011844E-6 |
| 8.420709512685392E-6 |
| 1.041448341728929... |
| 8.316075414862431E-6 |
| 9.721183814992126E-6 |
| 8.029436027707578E-6 |
| 6.307432259386365E-6 |
+-----+

```

only showing top 20 rows

Practical No 8

1. What day had the Peak high in price
2. What is the mean of the close column
3. What is the max and min of the volume column
4. How many days was the close lower than 220
5. What percentage of the time was the high greater than 250
6. What is the max high per year
7. What is the average close for each Calendar Month

```
from pyspark.sql import SparkSession
spark = SparkSession.builder.appName('walmart').getOrCreate()
df = spark.read.csv('walmart_stock.csv', inferSchema=True, header=True)
print(df.orderBy(df['High'].desc()).select(['Date']).head(1)[0]['Date'])
from pyspark.sql.functions import mean
df.select(mean('Close')).show()
from pyspark.sql.functions import min, max
df.select(max('Volume'), min('Volume')).show()
print(df.filter(df['Close'] < 220).count())
print(df.filter('High > 250').count() * 100/df.count())
from pyspark.sql.functions import (dayofmonth, hour,
dayofyear, month,
year, weekofyear,
format_number, date_format)
year_df = df.withColumn('Year', year(df['Date']))
year_df.groupBy('Year').max()['Year', 'max(High)'].show()
#Create a new column Month from existing Date column
month_df = df.withColumn('Month', month(df['Date']))
#Group by month and take average of all other columns
month_df = month_df.groupBy('Month').mean()
#Sort by month
month_df = month_df.orderBy('Month')
#Display only month and avg(Close), the desired columns
month_df['Month', 'avg(Close)'].show()
```

2015-01-13 00:00:00

avg(Close)
72.38844998012726

max(Volume)	min(Volume)
80898100	2094900

1258

0.0

Year	max(High)
2015	90.970001
2013	81.370003
2014	88.089996
2012	77.599998
2016	75.190002

Month	avg(Close)
1	71.44801958415842
2	71.306804443299
3	71.77794377570092
4	72.97361900952382
5	72.30971688679247
6	72.4953774245283
7	74.43971943925233
8	73.02981855454546
9	72.18411785294116
10	71.57854545454543
11	72.1110893069307
12	72.84792478301885

Practical No 9:

Spark SQL connecting with Data Source : Display all the High and Closing price

What is the average high price , close price, low price

What is the lowest price in high, close and low price

```
from pyspark.sql import SparkSession
spark = SparkSession.builder.appName('walmart').getOrCreate()
df = spark.read.csv('walmart_stock.csv', inferSchema=True, header=True)
print(df.orderBy(df['High'].desc()).select(['Date']).head(1)[0]['Date'])
from pyspark.sql.functions import mean

avg_high = df.agg({"High": "avg"}).collect()[0][0]
print("average high: " + str(avg_high))

avg_close = df.agg({"Close": "avg"}).collect()[0][0]
print("average close: " + str(avg_close))

avg_low = df.agg({"Low": "avg"}).collect()[0][0]
print("average low: " + str(avg_low))

low_high = df.agg({"High": "min"}).collect()[0][0]
print("low high: " + str(low_high))

low_close = df.agg({"Close": "min"}).collect()[0][0]
print("low close: " + str(low_close))

low_low = df.agg({"Low": "min"}).collect()[0][0]
print("low low: " + str(low_low))
```

```
2015-01-13 00:00:00
+-----+-----+-----+-----+-----+-----+-----+
| Date | Open | High | Low | Close | Volume | Adj Close |
+-----+-----+-----+-----+-----+-----+-----+
| 2012-01-03 00:00:00 | 59.970001 | 61.060001 | 59.869999 | 60.330002 | 12668800 | 52.619234999999996 |
| 2012-01-04 00:00:00 | 60.209998999999996 | 60.349998 | 59.470001 | 59.709998999999996 | 9593300 | 52.078475 |
| 2012-01-05 00:00:00 | 59.349998 | 59.619999 | 58.369999 | 59.419998 | 12768200 | 51.825539 |
| 2012-01-06 00:00:00 | 59.419998 | 59.450001 | 58.869999 | 59.0 | 8069400 | 51.45922 |
| 2012-01-09 00:00:00 | 59.029999 | 59.549999 | 58.919998 | 59.18 | 6679300 | 51.616215000000004 |
| 2012-01-10 00:00:00 | 59.43 | 59.709998999999996 | 58.98 | 59.040001000000004 | 6907300 | 51.494109 |
| 2012-01-11 00:00:00 | 59.060001 | 59.529999 | 59.040001000000004 | 59.400002 | 6365600 | 51.808098 |
| 2012-01-12 00:00:00 | 59.790001000000004 | 60.0 | 59.400002 | 59.5 | 7236400 | 51.895315999999994 |
| 2012-01-13 00:00:00 | 59.18 | 59.610001000000004 | 59.009997999999996 | 59.540001000000004 | 7729300 | 51.930203999999996 |
| 2012-01-17 00:00:00 | 59.869999 | 60.110001000000004 | 59.52 | 59.849998 | 8500000 | 52.200581 |
| 2012-01-18 00:00:00 | 59.790001000000004 | 60.029999 | 59.650002 | 60.009997999999996 | 5911400 | 52.340131 |
| 2012-01-19 00:00:00 | 59.93 | 60.73 | 59.75 | 60.610001000000004 | 9234600 | 52.863447 |
| 2012-01-20 00:00:00 | 60.75 | 61.25 | 60.669998 | 61.009997999999996 | 10378800 | 53.212320999999996 |
| 2012-01-23 00:00:00 | 60.810001 | 60.98 | 60.509997999999996 | 60.91 | 7134100 | 53.125104 |
| 2012-01-24 00:00:00 | 60.75 | 62.0 | 60.75 | 61.389998999999996 | 7362800 | 53.54375400000001 |
| 2012-01-25 00:00:00 | 61.18 | 61.610001000000004 | 61.040001000000004 | 61.470001 | 5915800 | 53.61353100000001 |
| 2012-01-26 00:00:00 | 61.799999 | 61.84 | 60.77 | 60.970001 | 7436200 | 53.177436 |
| 2012-01-27 00:00:00 | 60.860001000000004 | 61.119999 | 60.540001000000004 | 60.709998999999996 | 6287300 | 52.950665 |
| 2012-01-30 00:00:00 | 60.470001 | 61.32 | 60.349998 | 61.299999 | 7636900 | 53.465256999999994 |
| 2012-01-31 00:00:00 | 61.529999 | 61.57 | 60.580002 | 61.360001000000004 | 9761500 | 53.517590000000006 |
+-----+-----+-----+-----+-----+-----+-----+
only showing top 20 rows

average high: 72.83938807631165
average close: 72.38844998012726
average low: 71.9186009594594
low high: 57.060001
low close: 56.419998
low low: 56.299999
```

Practical No 10:

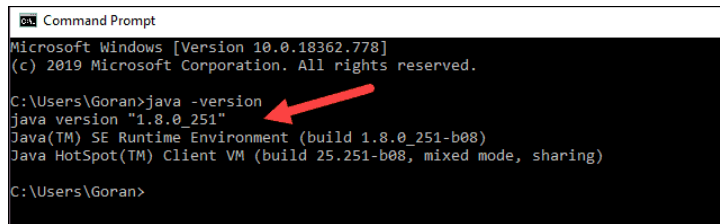
Installation SPark and Scala on window based application

Install Java 8

Apache Spark requires Java 8. You can check to see if Java is installed using the command prompt.

Open the command line by clicking Start > type cmd > click Command Prompt.

Type the following command in the command prompt:



```
Command Prompt
Microsoft Windows [Version 10.0.18362.778]
(c) 2019 Microsoft Corporation. All rights reserved.

C:\Users\Goran>java -version
java version "1.8.0_251"
Java(TM) SE Runtime Environment (build 1.8.0_251-b08)
Java HotSpot(TM) Client VM (build 25.251-b08, mixed mode, sharing)

C:\Users\Goran>
```

Install Apache Spark

Installing Apache Spark involves extracting the downloaded file to the desired location.

1. Create a new folder named Spark in the root of your C: drive. From a command line, enter the following:

```
cd \
```

```
mkdir Spark
```

2. In Explorer, locate the Spark file you downloaded.

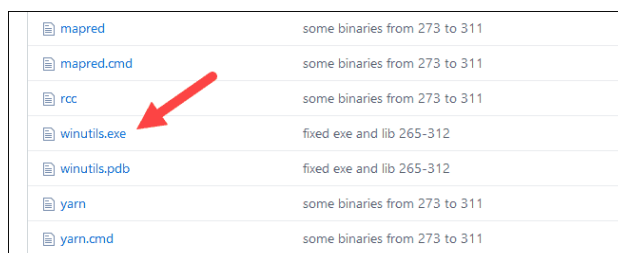
3. Right-click the file and extract it to C:\Spark using the tool you have on your system (e.g., 7-Zip).







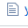
4. Now, your C:\Spark folder has a new folder spark-2.4.5-bin-hadoop2.7 with the necessary files inside.

Add winutils.exe File

Download the winutils.exe file for the underlying Hadoop version for the Spark installation you downloaded.

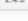
1. Navigate to this URL <https://github.com/cdarlint/winutils> and inside the bin folder, locate winutils.exe, and click it.



 mapred	some binaries from 273 to 311
 mapred.cmd	some binaries from 273 to 311
 rcc	some binaries from 273 to 311
 winutils.exe	fixed exe and lib 265-312
 winutils.pdb	fixed exe and lib 265-312
 yarn	some binaries from 273 to 311
 yarn.cmd	some binaries from 273 to 311

2. Find the Download button on the right side to download the file.

3. Now, create new folders Hadoop and bin on C: using Windows Explorer or the Command Prompt.


2.4.5

[Jobs](#)
[Stages](#)
[Storage](#)
[Environment](#)
[Executors](#)

Spark shell application

Executors

[Show Additional Metrics](#)

Summary

	RDD Blocks	Storage Memory	Disk Used	Cores	Active Tasks	Failed Tasks	Complete Tasks	Total Tasks
Active(1)	0	0.0 B / 434 MB	0.0 B	4	0	0	0	0
Dead(0)	0	0.0 B / 0.0 B	0.0 B	0	0	0	0	0
Total(1)	0	0.0 B / 434 MB	0.0 B	4	0	0	0	0

Executors

Show 20 entries

Search:

Executor ID	Address	Status	RDD Blocks	Storage Memory	Disk Used	Cores	Active Tasks	Failed Tasks	Complete Tasks
driver	DESKTOP-SFBGHOU:61547	Active	0	0.0 B / 434 MB	0.0 B	4	0	0	0

Showing 1 to 1 of 1 entries

[Previous](#)
1
[Next](#)